

# Introduction to Operating Systems

Operating System: Three Easy Pieces

---

# What happens when a program runs?

A running program executes instructions.

1. The processor **fetches** an instruction from memory.
2. **Decode**: Figure out which instruction this is
3. **Execute**: i.e., add two numbers, access memory, check a condition, jump to function, and so forth.
4. The processor moves on to the **next instruction** and so on.

This is also known as the **Von Neumann model** of computing.

Expectations for today's machines are much higher than this. For instance, multiple applications running at the same time.

# Operating System (OS)

Responsible for

- ◆ Making it easy to **run** programs
- ◆ Allowing programs to **share** memory
- ◆ Enabling programs to **interact** with devices

**OS is in charge of making sure the system operates  
correctly and efficiently.**

# Virtualization

The OS takes a **physical resource** and transforms it into a **virtual form** of itself.

- ◆ **Physical resource:** Processor, Memory, Disk ...
- ◆ The virtual form is more general, powerful and easy-to-use.
- ◆ Sometimes, we refer to the OS as a **virtual machine**.

# System call

So, how can we use all these nice features provided by an OS?

System calls allow users **to tell the OS what to do**.

- The OS provides some interfaces (APIs, standard library).
- A typical OS exports a few hundred system calls.
  - Run programs
  - Access memory
  - Access devices

# The OS is a resource manager.

- The OS **manages resources** such as *CPU*, *memory* and *disk*.
- The OS allows
  - Many programs to run → Sharing the CPU
  - Many programs to *concurrently* access their own instructions and data → Sharing memory
  - Many programs to access devices → For example, sharing disks

# Virtualizing the CPU

- A computer system has a certain number of physical/virtual CPUs.
  - For example, my laptop has 20 cores in a single processing unit.
  - A server machine can have tens of cores spread across several processing units
- Usually, the number of applications running on our machines is much larger than the number of available CPUs or cores
- The OS turns a single CPU into a seemingly infinite number of CPUs (it's just an illusion!).
  - Allowing many programs to seemingly run at once → **Virtualizing the CPU**

# Virtualizing the CPU (Cont.)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <sys/time.h>
4      #include <assert.h>
5      #include "common.h"
6
7      int
8      main(int argc, char *argv[])
9      {
10         if (argc != 2) {
11             fprintf(stderr, "usage: cpu <string>\n");
12             exit(1);
13         }
14         char *str = argv[1];
15         while (1) {
16             Spin(1); // Repeatedly checks the time and
17                      // returns once it has run for a second
18             printf("%s\n", str);
19         }
20         return 0;
21     }
```

**Simple Example(cpu.c): Code That Loops and Prints**



# Virtualizing the CPU (Cont.)

- Execution result 1.

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

**Run forever; Only by pressing "Control-c" can we halt the program**

# Virtualizing the CPU (Cont.)

- Execution result 2.

```
prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...
```

**Even though we have only **one processor**, all four of programs seem to be running **at the same time**!**

# Virtualizing Memory

- The physical memory is *an array of bytes*.
- A program keeps all of its data structures in memory.
  - **Read memory** (load):
    - Specify an address to be able to access the data
  - **Write memory** (store):
    - Specify the data to be written to the given address

# Virtualizing Memory (Cont.)

- A program that Accesses Memory (mem.c)

```
1      #include <unistd.h>
2      #include <stdio.h>
3      #include <stdlib.h>
4      #include "common.h"
5
6      int
7      main(int argc, char *argv[])
8      {
9          int *p = malloc(sizeof(int)); // a1: allocate some
                                         memory
10         assert(p != NULL);
11         printf("(d) address of p: %08x\n",
12              getpid(), (unsigned) p); // a2: print out the
                                         address of the memory
13         *p = 0; // a3: put zero into the first slot of the memory
14         while (1) {
15             Spin(1);
16             *p = *p + 1;
17             printf("(d) p: %d\n", getpid(), *p); // a4
18         }
19         return 0;
20     }
```

# Virtualizing Memory (Cont.)

- The output of the program `mem.c`

```
prompt> ./mem 0
(2134) memory address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- The newly allocated memory is at address `00200000`.
- It updates the value and prints out the result.

# Virtualizing Memory (Cont.)

- Set to 0 /proc/sys/kernel/randomize\_va\_space
- Running `mem.c` multiple times

```
prompt> ./mem 0 & ./mem 0 &  
[1] 24113  
[2] 24114  
(24113) memory address of p: 00200000  
(24114) memory address of p: 00200000  
(24113) p: 1  
(24114) p: 1  
(24114) p: 2  
(24113) p: 2  
(24113) p: 3  
(24114) p: 3  
...
```

- It is as if each running program has its **own private memory**.
  - Each running program has allocated memory at the same address.
  - Each seems to be updating the value at `00200000` independently.

# Virtualizing Memory (Cont.)

- Each process accesses its own private **virtual address space**.
  - The OS maps **address space** onto the **physical memory**.
  - A memory reference within one running program does not affect the address space of other processes.
  - Physical memory is a shared resource, managed by the OS.

# The problem of Concurrency

- The OS is juggling **many things at once**, first running one process, then another, and so forth.
- Modern **multi-threaded programs** also exhibit the concurrency problem.



# Concurrency Example (Cont.)

```
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads
21                             <value>\n");
22             exit(1);
23         }
24         loops = atoi(argv[1]);
25         pthread_t p1, p2;
26         printf("Initial value : %d\n", counter);
27         Pthread_create(&p1, NULL, worker, NULL);
28         Pthread_create(&p2, NULL, worker, NULL);
29         Pthread_join(p1, NULL);
30         Pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

- The main program creates **two threads**.
  - Thread: a function running within the same memory space. Each thread starts running in a routine called `worker()`.
  - `worker()`: increments a counter

# Concurrency Example

## A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include "common.h"
4
5      int counter = 0;
6      int loops;
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
15     ...
```

# Concurrency Example (Cont.)

- `loops` determines how many times each of the two workers will **increment the shared counter** in a loop.

- `loops: 1000.`

```
prompt> gcc -o threads threads.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- `loops: 100000.`

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

# How many instructions are in *counter++*?

## A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include "common.h"
4
5      int counter = 0;
6      int loops;
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
15     ...
```

# Why is this happening?

- Increment a shared counter → takes three instructions.
  - Load the value of the counter from memory into register.
  - Increment it
  - Store it back into memory
- These three instructions do not execute **atomically**. → **Concurrency** problem.

# Persistence

- Devices such as DRAM store values in a volatile fashion.
- *Hardware* and *software* are needed to store data **persistently**.
  - **Hardware:** I/O device such as a hard drive or solid-state drives (SSDs)
  - **Software:**
    - File system manages the disk.
    - File system is responsible for storing any files the user creates.

# Persistence (Cont.)

- Create a file (`/tmp/file`) that contains the string "hello world"

```
1      #include <stdio.h>
2      #include <unistd.h>
3      #include <assert.h>
4      #include <fcntl.h>
5      #include <sys/types.h>
6
7      int
8      main(int argc, char *argv[])
9      {
10         int fd = open("/tmp/file", O_WRONLY | O_CREAT
                        | O_TRUNC, S_IRWXU);
11         assert(fd > -1);
12         int rc = write(fd, "hello world\n", 13);
13         assert(rc == 13);
14         close(fd);
15         return 0;
16     }
```

`open()`, `write()`, and `close()` system calls are routed to the part of OS called the file system, which handles the requests

# Persistence (Cont.)

- What does the OS do in order to write to disk?
  - Deals with additional abstraction layers, such as RAID, Logical Volume Manager, and so forth
  - Figures out **where** on disk this new data will reside
  - **Issue I/O** requests to the underlying storage device
- File system handles system crashes during write.
  - **Journaling** or **copy-on-write**
  - Carefully ordering writes to disk



# Design Goals

- Build up **abstraction**
  - Make the system convenient and easy to use.
- Provide high **performance**
  - Minimize the overhead of the OS.
  - OS must strive to provide virtualization without excessive overhead.
- **Protection** between applications
  - Isolation: Bad behavior of one does not harm other and the OS itself.

# Design Goals (Cont.)

- High degree of **reliability**
  - The OS must also run non-stop.
- Other issues
  - Energy-efficiency
  - Security
  - Mobility