

# 15. Address Translation

Operating System: Three Easy Pieces

---

# Memory Virtualizing with Efficiency and Control

- Why memory virtualization?
  - Two main reasons: Flexibility and Access Limitation.
- In Limited Direct Execution (LDE) for CPU virtualization, we used some hardware support to realize an efficient mechanism so that the OS could always be in control
- We want to obtain the same results (efficiency+control) in memory virtualization
  - Similarly to CPU virtualization, we will need some support from the hardware
  - Very simple at the beginning, but will grow fairly complex then.
- We want to study here mechanisms to realize an efficient virtualization of memory.

# Address Translation

- The hardware support consists in the transformation of a **virtual address** to a **physical address**.
  - The desired information is actually stored in a physical address.
- The hardware, by itself, only provides mechanisms for the translation.
  - The OS is the entity that drives the hardware in the translation process
- Just like CPU virtualization, we want to create an **illusion**: the illusion that every process has all the memory for itself

# 3 Assumptions

- We will start the analysis of memory virtualization with 3 (unrealistic) assumptions:
  - The address space of a process must be placed contiguously in physical memory
  - The size of the address space is smaller than the size of the physical memory
  - Each address space has the same size
- Let's start with something simple, then we will relax these assumptions.

# Example: Application Accessing Memory


```
void func() {  
    int x;  
    ...  
    x = x + 3; // this is the line of code we are interested in  
}
```

- In C code:
  - **Load** a value from memory
  - **Increment** it by three
  - **Store** the value back into memory

```
128 : movl 0x0(%ebx), %eax      ; load 0+ebx into eax  
132 : addl $0x03, %eax         ; add 3 to eax register  
136 : movl %eax, 0x0(%ebx)     ; store eax back to mem
```

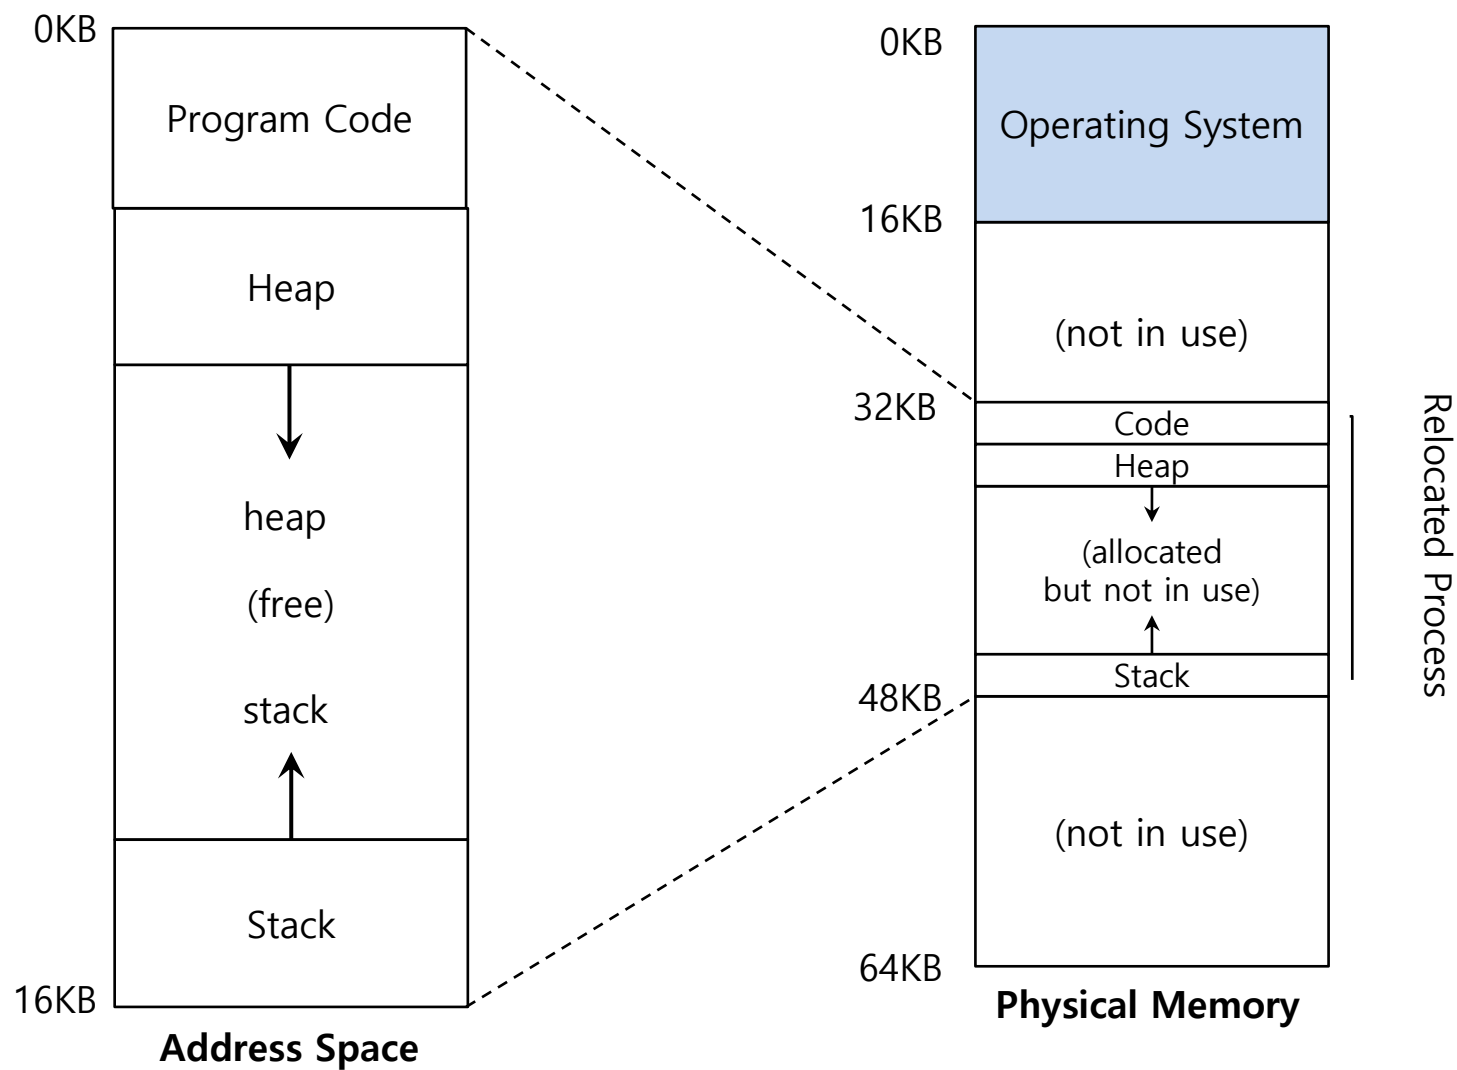
- In Assembly code:
  - **Load** the value at that address into `eax` register.
  - **Add** 3 to `eax` register.
  - **Store** the value in `eax` back into memory.

# Example: Application Accessing Memory (Cont.)

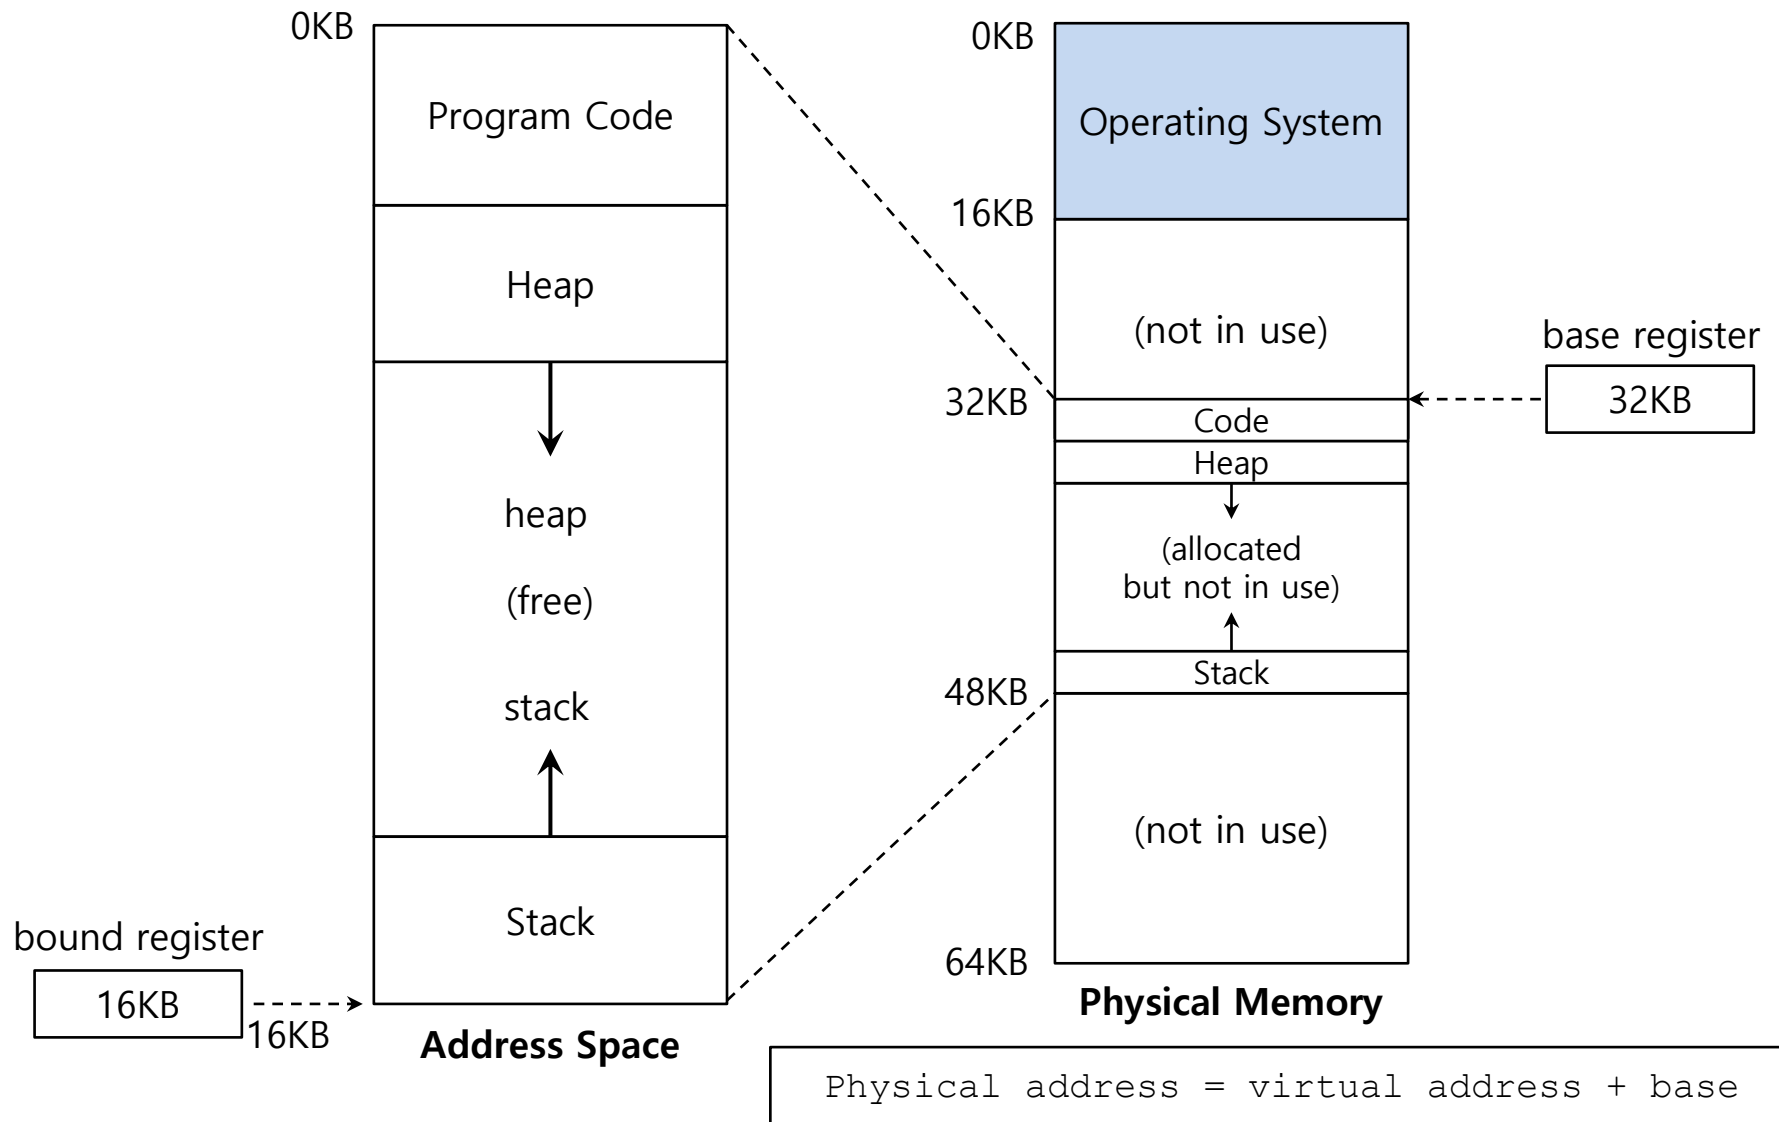


- Fetch instruction at address 128
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 136
- Execute this instruction (store to address 15 KB)

# A Single Relocated Process



# Base and Bound Registers





# Relocation and Address Translation

128 : `movl 0x0(%ebx), %eax`

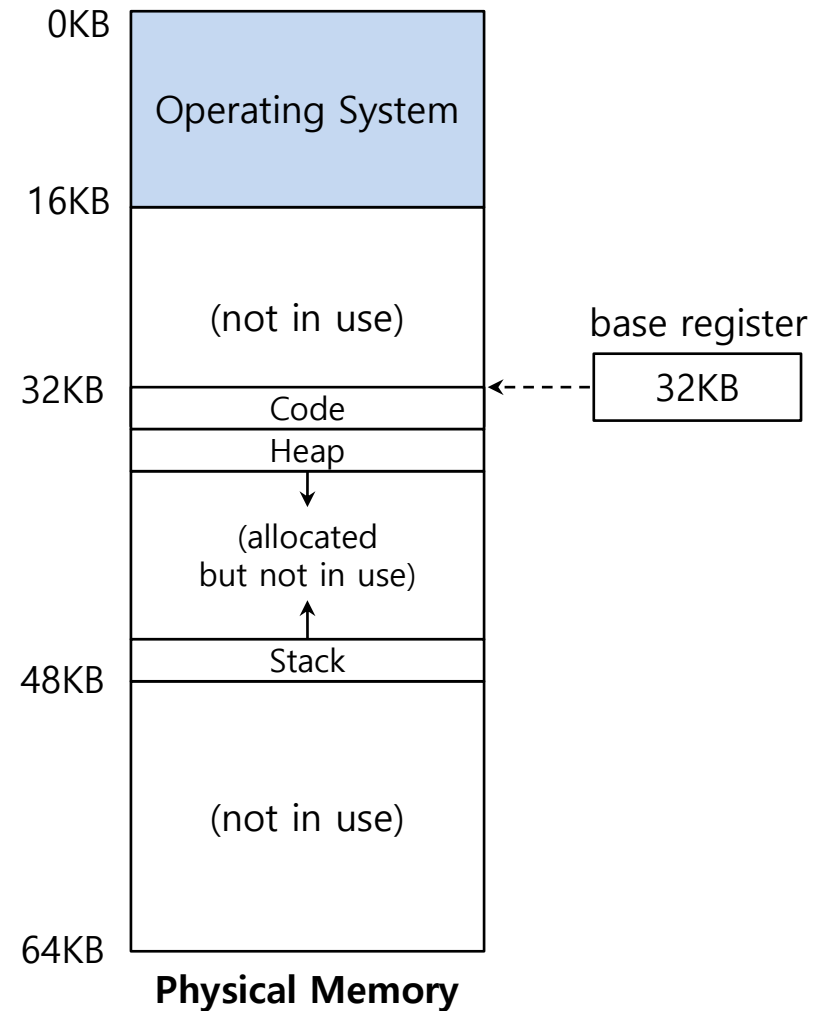
- **Fetch** instruction at address 128

$128 \rightarrow 128 + 32768 = 32896$

- **Execute** this instruction

- Load from address 15KB

$15K \rightarrow 15K + 32K = 47K$



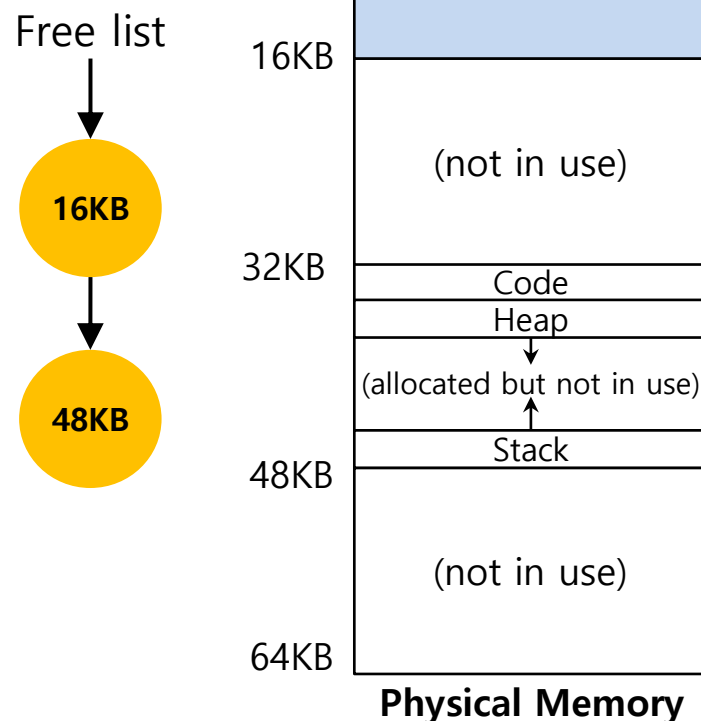
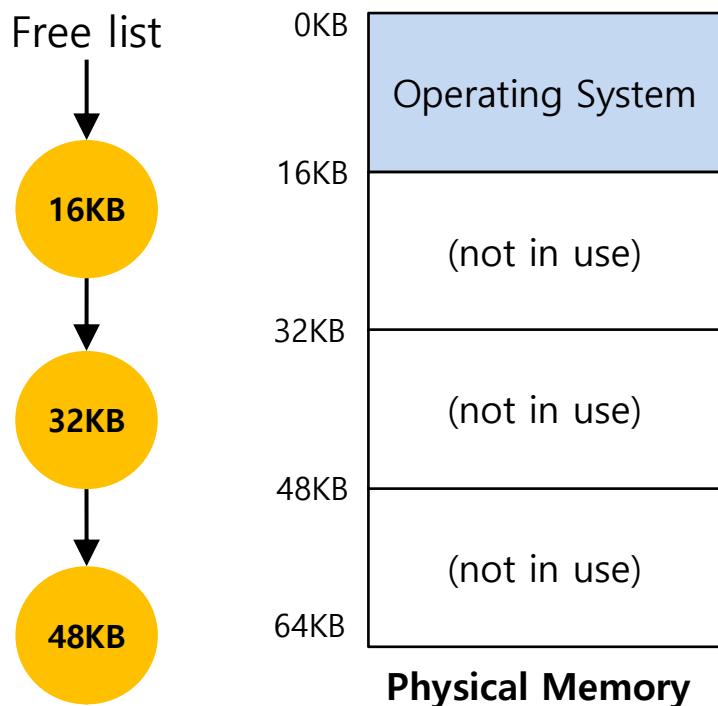
# OS Issues for Memory Virtualizing

- Hardware support is good, but we also need some OS implementation
- The OS must **take action** to implement **base-and-bounds** approach.
- Three critical junctures:
  - When a process **starts running**:
    - Finding space for address space in physical memory
  - When a process is **terminated**:
    - Reclaiming the memory for use
  - When context **switch occurs**:
    - Saving and storing the base-and-bounds pair

# OS Issues: When a Process Starts Running



- The OS must **find room** for a new address space.
  - free list : A list of the ranges of the physical memory which are not in use.
  - Easy with 16KB slots

The OS searches the free list



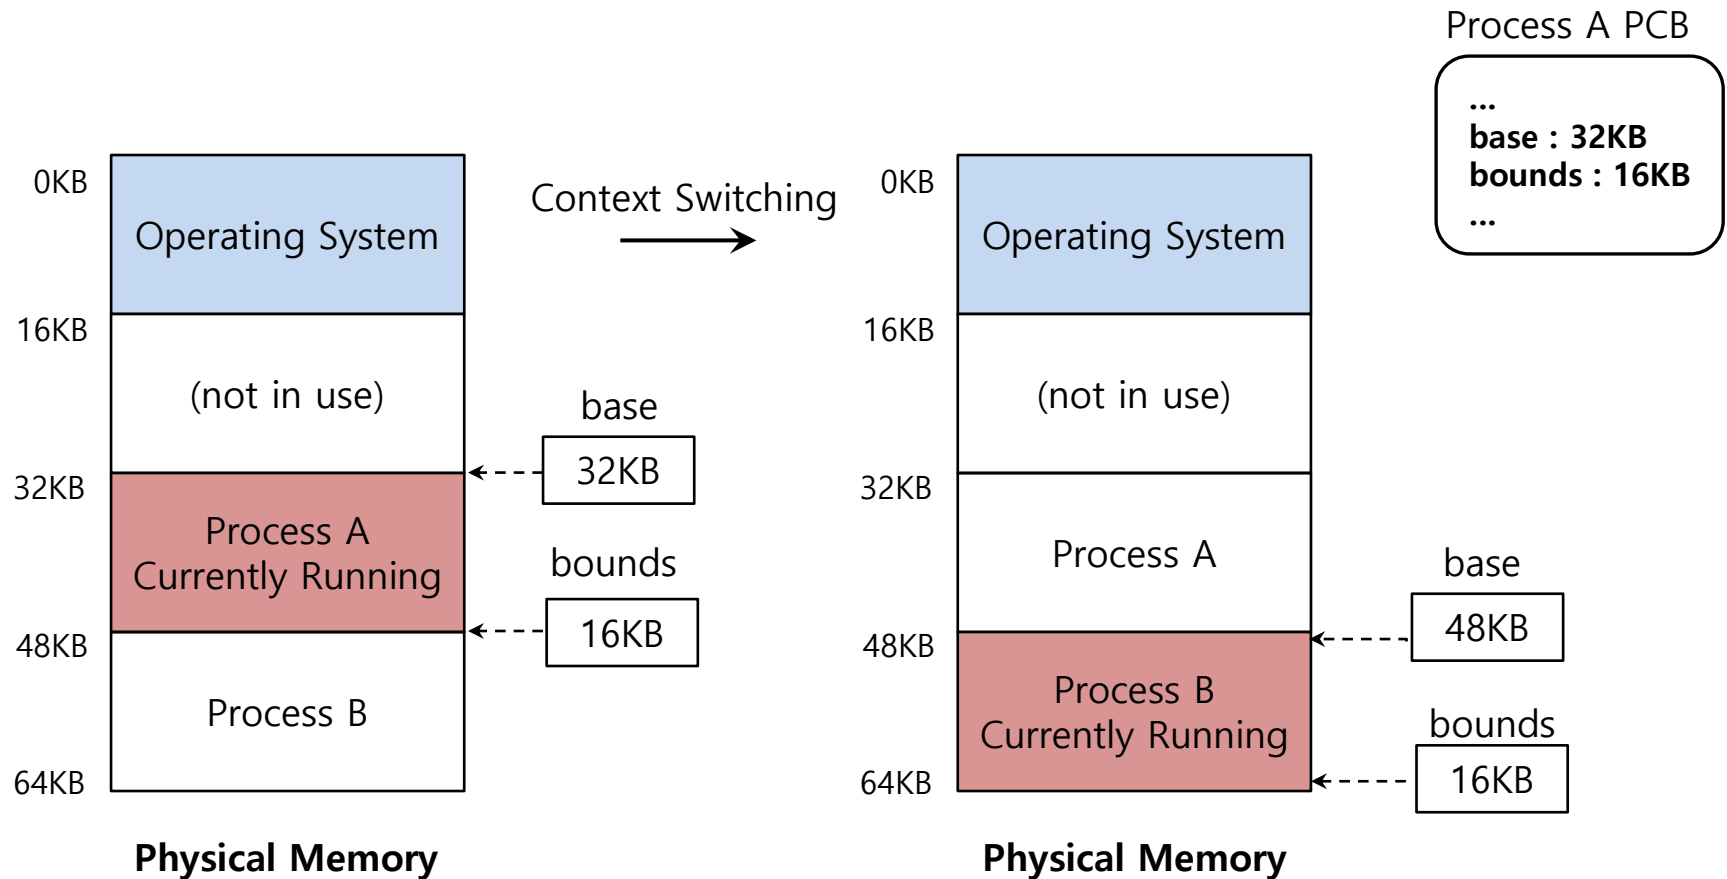
# OS Issues: When a Process Is Terminated

- The OS must **put the memory back** on the free list.



# OS Issues: When Context Switch Occurs

- The OS must **save and restore** the base-and-bounds pair.
  - In **process structure** or **process control block (PCB)**



# Summary of Hardware Requirements so far

| Hardware Requirements   | Notes  |
|---|--|
| Privileged mode   | Needed to prevent user-mode processes from executing privileged operations     |
| Base/bound registers  | Need pair of registers per CPU to support address translation and bound checks |
| Ability to translate virtual addresses and check if within bounds | Circuitry to make translations and check limits, in this case quite simple     |
| Privileged instruction(s) to update base/bounds                   | OS must be able to set these values before letting a user program run          |
| Privileged instruction(s) to register exception handlers          | OS must be able to tell hardware what code to run if exception occurs          |
| Ability to raise exceptions                                       | When processes try to access privileged instructions or out-of-bounds memory   |